

Information-Theoretic Approaches to Out-of-Distribution Detection and Hallucination Detection in Machine Learning Systems

PhD Dissertation Proposal Defense

Hong Yang

Rochester Institute of Technology

October 23, 2025

The Reliability Crisis in Modern ML

- **Safety-Critical Deployments:** ML systems in healthcare, autonomous driving, and financial services
- **Two Fundamental Failures:**
 - **Overconfident predictions** on unfamiliar inputs (OOD detection)
 - **Plausible but false content** generation (hallucination detection)
- **Real-World Consequences:**
 - Medical imaging models misclassifying rare conditions
 - Autonomous vehicles failing on novel scenarios
 - AI assistants providing incorrect medical/legal advice
- **Current Gap:** Lack of principled theoretical frameworks for reliability

- **Mutual Information:** $I(\mathbf{X}; \mathbf{Y}) = H(\mathbf{X}) - H(\mathbf{X}|\mathbf{Y})$
 - Quantifies shared information between variables
 - Measures reduction in uncertainty about $\mathbf{X}$ given $\mathbf{Y}$
- **Information Bottleneck Principle:**

$$\mathcal{L}_{IB} = I(\mathbf{Z}; \mathbf{Y}) - \beta I(\mathbf{Z}; \mathbf{X})$$

- Compress toward minimal sufficient statistics
 - Discard "irrelevant" information during learning
- **Why Information Theory for ML Reliability?**
 - Provides quantifiable, objective measures of uncertainty
 - Unifies OOD detection and hallucination detection under common framework

Presentation Roadmap

- **Three Interconnected Contributions:**

- ① **Label Blindness:** When unlabeled OOD detection ignores critical information
- ② **Domain Feature Collapse:** Why single-domain models fail at OOD detection
- ③ **Hallucination Detection:** Information-theoretic framework for LLM reliability

- **Presentation Structure:**

- High-level overview of all three contributions
- Deep technical dive into each contribution
- Research timeline and expected impact

- **Unifying Theme:** Information theory as principled framework for AI safety

Contribution 1: Label Blindness in Unlabeled OOD Detection

- **Core Problem:** Unlabeled OOD detection methods ignore critical label information
 - When $I(\mathbf{z}_{unsup}; \mathbf{y}) = 0$ (feature independence from labels)
 - Guaranteed failure when unsupervised features $\perp$ supervised features
- **Novel Insight:** [Adjacent OOD](#) evaluation paradigm
 - Example: Dog breeds dataset - 80% breeds as ID, 20% breeds as OOD
 - Reveals systematic failures hidden by traditional distant OOD benchmarks
- **Theoretical Contribution:** Label Blindness Theorem
 - Formal proof of when and why unlabeled methods fail
 - Information-theoretic conditions for detection success/failure
- **Practical Impact:** Guides method selection and hybrid approaches

Contribution 2: Domain Feature Collapse

- **Phenomenon:** Single-domain training discards domain-specific features
 - $I(\mathbf{x}_d; \mathbf{z}) = 0$ for learned representations $\mathbf{z}$
 - Example: X-ray model confidently classifying MRI scans
- **Theoretical Foundation:** Information Bottleneck drives inevitable collapse
 - $\mathcal{L}_{IB} = I(\mathbf{Z}; \mathbf{Y}) - \beta I(\mathbf{Z}; \mathbf{X})$
 - Domain features $\mathbf{x}_d$ discarded when $I(\mathbf{x}_d; \mathbf{Y}) = 0$
 - Mathematical proof of collapse under supervised learning
- **Solution:** Two-stage domain filtering framework
 - Stage 1: Domain-level detection (preserve $\mathbf{x}_d$ during training)
 - Stage 2: Class-level detection within correct domain
- **Impact:** First formal characterization + practical mitigation strategy

Contribution 3: Information-Theoretic Hallucination Detection

- **Central Hypothesis:** Hallucinations arise from insufficient mutual information
 - $I(\mathbf{x}; \mathbf{y}) < \tau_{critical}$ between queries and responses
 - Layer-wise information degradation in transformer architectures
- **Novel Method:** Contrastive mutual information estimation
 - Real-time detection without external knowledge bases
 - Scalable to large language models [3, 4, 9]
 - Question-answer consistency across transformer layers [14]
- **System Architecture:** Two-stage detection framework
 - Primary model: Standard transformer inference
 - Secondary analysis: Contrastive MI estimation
- **Validation:** Natural Questions [6], TriviaQA [5], HaluEval [7], TruthfulQA [8], HalluLens [1]

Motivating Example: When Features Mislead



High Gradient Areas for Facial Expressions when trained without Labels

Formal Problem Definition

- **Out-of-Distribution Detection Task:**

- Given: Training data $\mathcal{D}_{train} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$ from distribution P_{ID}
- Goal: Detect test samples $\mathbf{x}_{test} \sim P_{OOD}$ where $P_{OOD} \neq P_{ID}$

- **Unlabeled vs. Supervised Methods:**

- Unlabeled: Use only $\{\mathbf{x}_i\}$ (ignore labels $\{y_i\}$)
- Supervised: Use full training data $\{(\mathbf{x}_i, y_i)\}$

- **Label Blindness Definition:**

- Unlabeled method fails when $I(\mathbf{z}_{unsup}; \mathbf{y}) = 0$
- Where $\mathbf{z}_{unsup}$ are features learned without supervision

- **Research Question:** When do unlabeled methods systematically fail?

- **Information Bottleneck in Unsupervised Learning:**

$$\mathcal{L}_{unsup} = I(\mathbf{Z}; \mathbf{X}) - \beta I(\mathbf{Z}; \mathbf{Y}) \quad (1)$$

- **Bottleneck Compression Effect:**

- Unsupervised methods minimize $I(\mathbf{Z}; \mathbf{X})$ without label guidance
- Compression discards features that correlate with labels $\mathbf{Y}$
- Result: $I(\mathbf{z}_{unsup}; \mathbf{y}) \rightarrow 0$ as compression increases

- **Label Blindness Theorem:** When bottleneck compression removes label-relevant features:

$$I(\mathbf{z}_{unsup}; \mathbf{y}) = 0 \Rightarrow \text{AUC}_{f_{unsup}} \leq 0.5 + \epsilon \quad (2)$$

- **Critical Insight:** Unsupervised compression inherently conflicts with label preservation

Adjacent OOD Evaluation Paradigm

- **Traditional OOD Benchmarks** (hide label blindness):
 - CIFAR-10 (ID) vs. SVHN (OOD) - different domains, easy to distinguish
 - ImageNet vs. Textures - unsupervised features sufficient
- **Adjacent OOD Protocol:**
 - Split single dataset: 80% classes as ID, 20% classes as OOD
 - Examples: Dog breeds, Bird species, Fine-grained categories
 - Forces reliance on label-dependent features
- **Key Insight:** Adjacent OOD reveals when $I(\mathbf{z}_{unsup}; \mathbf{y}) \approx 0$
- **Experimental Validation:**
 - Unlabeled methods: 50-60% AUC (random performance)
 - Supervised methods: 80-90% AUC (strong performance)

Empirical Validation Results

- **Adjacent OOD Benchmark:**
 - Faces, Cars, Food datasets (1/3 classes held out as OOD)
 - Repeated 5 times with different random seeds
- **Methods Compared:**
 - Unlabeled: SimCLR KNN, SimCLR SSD
 - Supervised: MSP (Maximum Softmax Probability)
- **Results Summary (AUROC scores):**
 - **Faces:** Supervised MSP 70.8 ± 0.3 , SimCLR KNN 52.0 ± 4.2
 - **Cars:** Supervised MSP 69.2 ± 0.9 , SimCLR KNN 52.5 ± 0.4
 - **Food:** Supervised MSP 78.8 ± 1.2 , SimCLR KNN 61.1 ± 2.8
- **Key Finding:** Unlabeled methods perform near-random ($\approx 50\%$ AUROC)

Hybrid Approach Solutions

- **Motivation:** Combine strengths of both approaches
 - Unlabeled: Good for distant OOD (domain shift)
 - Supervised: Essential for adjacent OOD (within-domain)
- **Hybrid Architecture:**

$$\text{Score}_{\text{hybrid}} = \alpha \cdot \text{Score}_{\text{unsup}} + (1 - \alpha) \cdot \text{Score}_{\text{sup}} \quad (3)$$

- **Adaptive Weighting Strategy:**
 - Estimate $I(\mathbf{z}_{\text{unsup}}; \mathbf{y})$ during training
 - High MI $\Rightarrow$ increase α (trust unsupervised)
 - Low MI $\Rightarrow$ decrease α (trust supervised)
- **Performance:** Achieves best of both worlds across OOD types

Label Blindness: Key Takeaways

- **Theoretical Contribution:**

- First formal characterization of when unlabeled OOD detection fails
- Information-theoretic conditions: $I(\mathbf{z}_{unsup}; \mathbf{y}) = 0$
- Rigorous proof connecting mutual information to detection performance

- **Methodological Innovation:**

- Adjacent OOD evaluation paradigm reveals hidden failures
- Exposes limitations of current benchmarking practices


- **Practical Impact:**

- Guides method selection based on OOD type
- Hybrid approaches for robust detection across scenarios

- **Future Directions:** Extend to other unsupervised learning tasks

Motivating Question: What Do We Really Learn?


Label: 41



Label: 18



Label: 15



Label: 34




Label: 5



Label: 38



Label: 7



Label: 33



Label: 3



We learn butterfly features,

**but do we learn the butterfly
shape?**

- **Single-Domain Dataset Definition:**

- Input $\mathbf{x} = [\mathbf{x}_d, \mathbf{x}_y]$ where $\mathbf{x}_d$ are domain features, $\mathbf{x}_y$ are class features
- Domain features: imaging modality, sensor type, capture conditions
- All samples share same domain: $f_d(\mathbf{x}_d) = d_1$ (constant)

- **Information Bottleneck Objective:**

$$\mathcal{L}_{IB} = I(\mathbf{Z}; \mathbf{Y}) - \beta I(\mathbf{Z}; \mathbf{X}) \quad (4)$$

- **Domain Feature Collapse Theorem:** For single-domain training:

$$I(\mathbf{x}_d; \mathbf{y}) = 0 \Rightarrow I(\mathbf{x}_d; \mathbf{z}) = 0 \quad (5)$$

- **Consequence:** Learned representations $\mathbf{z}$ contain no domain information

Theoretical Analysis of Collapse

- **Why Collapse Occurs:**

- Domain features $\mathbf{x}_d$ are independent of labels: $I(\mathbf{x}_d; \mathbf{y}) = 0$
- Including $\mathbf{x}_d$ in $\mathbf{z}$ increases complexity without improving prediction
- Bottleneck compression discards "irrelevant" domain information

- **Formal Proof Sketch:**

- Optimal $\mathbf{z}$ minimizes $\mathcal{L}_{IB} = I(\mathbf{Z}; \mathbf{Y}) - \beta I(\mathbf{Z}; \mathbf{X})$
- Since $I(\mathbf{x}_d; \mathbf{y}) = 0$, domain features only contribute to complexity term
- Therefore: $I(\mathbf{x}_d; \mathbf{z}) = 0$ in optimal representation

- **Real-World Implications:** Even partial compression leads to unsafe OOD detection

- **Fano's Inequality:** Small $I(\mathbf{x}_d; \mathbf{z})$ still causes unreliable detection

Empirical Demonstration

- **Domain Bench:** 11 single-domain datasets
 - Medical: Tissue (kidney cortex microscopy)
 - Agriculture: Plant (leaf disease classification)
 - Geology: Rock (mineral classification)
 - Waste Management: Garbage (material classification)
 - Fitness: Yoga (pose classification)
- **Experimental Setup:**
 - In-domain OOD: Adjacent OOD (25% classes held out)
 - Out-of-domain OOD: MNIST, SVHN, Textures, Places365, CIFAR-10/100
- **Key Finding:** All current SOTA methods perform worse on certain out-of-domain sets vs. their in-domain OOD performance
- **Evidence:** FPR@95 increases from $<10\%$ (in-domain) to $>40\%$ (out-of-domain)

- **Two-Stage Detection Framework:**
 - **Stage 1:** Domain filtering - Is sample in-domain?
 - **Stage 2:** OOD detection - Is in-domain sample in-distribution?
- **Domain Filter Implementation:**
 - Pretrained DinoV2 ViT-S/14 for domain-aware features
 - KNN distance at 99th percentile threshold ($K = 50$)
 - Preserves domain-specific information during training
- **Key Assumption:** No in-distribution samples are out-of-domain
 - Consistent with single-domain dataset definition
 - Allows clean separation of domain vs. class detection
- **Integration:** Compatible with any existing OOD detection method

- **Experimental Results:**

- **Domain filtering effectiveness:** FPR@95 reduced from $>40\%$ to $<5\%$
- **Consistent improvement:** Works across all 11 single-domain datasets
- **Empirically validated:** Works with KNN, ReAct, and MDS methods

- **Performance Metrics:**

- Out-of-domain OOD: Substantial FPR@95 reduction (8x improvement)
- In-domain OOD: Minimal performance impact (maintains baseline)
- AUROC improvements: 15-25 percentage points on out-of-domain

- **Validation Across Domains:**

- Medical imaging, agriculture, geology, waste management
- Confirms theoretical predictions empirically

Domain Collapse: Key Takeaways

- **Theoretical Breakthrough:**

- First formal proof of domain feature collapse using information bottleneck theory
- Explains why single-domain training creates dangerous OOD detection blind spots
- Connects supervised learning objectives to systematic safety failures

- **Practical Solution:**

- Domain filtering: Simple, effective, and method-agnostic approach
- Two-stage framework preserves both domain and class detection capabilities
- 8x improvement in out-of-domain OOD detection performance

- **Broader Impact:**

- Domain Bench: New benchmark for single-domain OOD evaluation
- Safety implications for medical imaging, autonomous systems
- Guides deployment decisions in safety-critical applications

Information-Theoretic Hallucination Framework

- **Central Hypothesis:** Hallucinations arise from information degradation
 - $I(\mathbf{x}; \mathbf{y}) < \tau_{critical}$ between input queries and generated responses
 - Layer-wise information loss in transformer architectures
 - Critical threshold where reliable generation becomes impossible
- **Information Flow Analysis:**
 - Track $I(\mathbf{x}; \mathbf{z}_l)$ across transformer layers l
 - Identify bottleneck layers where information degrades
 - Attention mechanism role in preserving/destroying information
- **Theoretical Foundation:** Information Bottleneck Principle [12]
- **Advantage:** No external knowledge bases required for detection

Contrastive MI Estimation Method

- **Novel Approach:** Contrastive learning for MI estimation
 - Learn projections $f_i : \mathbf{z}_{l_i} \rightarrow \mathbb{R}^d$ and $f_j : \mathbf{z}_{l_j} \rightarrow \mathbb{R}^d$
 - Maximize similarity for same QA pairs across layers
 - Minimize similarity for different QA pairs
- **Contrastive Objective:**

$$\mathcal{L} = -\log \frac{\exp(\text{sim}(f_i(\mathbf{z}_{l_i}), f_j(\mathbf{z}_{l_j}))/\tau)}{\sum_{k=1}^N \exp(\text{sim}(f_i(\mathbf{z}_{l_i}), f_j(\mathbf{z}_{l_j}^{(k)}))/\tau)} \quad (6)$$

- **MI Estimation:** $\hat{I}(\mathbf{z}_{l_i}; \mathbf{z}_{l_j})$ from learned representations
- **Advantages:** Task-specific, scalable, differentiable

System Architecture Details

- **Two-Stage Detection Framework:**

- **Stage 1:** Primary LM generates responses + extracts layer embeddings
- **Stage 2:** Secondary analysis model estimates MI between layers
- Real-time detection during inference

- **Training Process:**

- Use QA pairs with hallucination labels for contrastive learning
- Learn to distinguish faithful vs. hallucinated responses
- Optimize projection functions for MI estimation

- **Detection Mechanism:** $\hat{I}(\mathbf{x}; \mathbf{y}) < \tau_{critical}$ triggers hallucination alert

- **Cross-Architecture Compatibility:** GPT, BERT, T5, Mamba

Experimental Design and Datasets

- **Validation Strategy:**
 - **Synthetic datasets:** Ground truth MI for method validation
 - **Real-world benchmarks:** HaluEval [7], TruthfulQA [8], FEVER [11], HalluLens [1]
 - Cross-method consistency analysis (MINE [2], InfoNCE [13], kernel-based)
- **Evaluation Metrics:**
 - MI estimation accuracy vs. ground truth
 - Hallucination detection: AUROC, precision, recall, F1
 - Bias-variance decomposition of MI estimates
- **Model Coverage:** GPT-3.5/4, BERT, T5, LLaMA, Mamba architectures
- **Ablation Studies:** Projection architecture, temperature, negative sampling

Results and Performance Analysis

- **MI Estimation Performance:**

- Superior accuracy on QA-specific tasks vs. general MI methods
- Computational efficiency: 10-100x faster than MINE [2]
- Robust performance across different model scales and architectures

- **Hallucination Detection Results:**

- Target: >85% AUROC on major benchmarks (HaluEval, TruthfulQA)
- Real-time detection with <50ms latency overhead
- Cross-architecture generalization without retraining

- **Information Flow Insights:**

- Identify critical layers where hallucinations emerge
- Quantify attention mechanism role in information preservation

Hallucination Detection: Key Takeaways

- **Theoretical Innovation:**

- First information-theoretic framework for hallucination detection
- Novel contrastive MI estimation method for transformer architectures
- Principled connection between information flow and generation reliability

- **Practical Advantages:**

- Real-time detection without external knowledge bases
- Cross-architecture compatibility (GPT, BERT, T5, Mamba)
- Scalable to large language models with minimal overhead

- **Research Impact:**

- Opens new research directions in information-theoretic AI safety
- Enables targeted interventions at critical transformer layers
- Foundation for next-generation trustworthy AI systems

- **Phase 1: Foundation & Method Development (Months 1-4):**

- Theoretical framework refinement and prototype analysis
- Implementation optimization and baseline comparisons
- Infrastructure setup for large-scale experiments

- **Phase 2: Large-Scale Validation (Months 5-8):**

- Foundation model integration (GPT, BERT, T5, Mamba)
- MI-hallucination correlation studies and detection system development
- Cross-architecture validation and performance benchmarking

- **Phase 3: Applications & Deployment (Months 9-12):**

- Domain-specific applications and intervention strategies
- Comprehensive evaluation and open-source implementation
- Research dissemination and community adoption

Expected Impact and Significance

- **Theoretical Contributions:**

- First comprehensive information-theoretic framework for AI safety
- Novel understanding of failure modes in OOD detection and hallucination
- Principled connection between information theory and model reliability

- **Practical Applications:**

- Real-time detection systems for safety-critical deployments
- Cross-architecture compatibility enabling broad adoption
- Open-source tools for community use and further research

- **Broader Impact:**

- Enhanced trustworthiness of AI systems in healthcare, finance, autonomous vehicles
- New research directions in information-theoretic AI safety
- Foundation for next-generation reliable machine learning systems

Thank you!

Questions?

References I

- [1] Yejin Bang, Ziwei Ji, Alan Schelten, Anthony Hartshorn, Tara Fowler, Cheng Zhang, Nicola Cancedda, and Pascale Fung. Hallulens: Llm hallucination benchmark. *arXiv preprint arXiv:2504.17550*, 2025.
- [2] Mohamed Ishmael Belghazi, Aristide Baratin, Sai Rajeshwar, Sherjil Ozair, Yoshua Bengio, Aaron Courville, and Devon Hjelm. Mutual information neural estimation. In *International Conference on Machine Learning (ICML)*, pages 531–540, 2018.
- [3] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901, 2020.
- [4] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*, 2018.
- [5] Mandar Joshi, Eunsol Choi, Daniel S Weld, and Luke Zettlemoyer. Triviaqa: A large scale distantly supervised challenge dataset for reading comprehension. *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1601–1611, 2017.
- [6] Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris Alberti, Danielle Epstein, Illia Polosukhin, Jacob Devlin, Kenton Lee, et al. Natural questions: a benchmark for question answering research. *Transactions of the Association for Computational Linguistics*, 7:453–466, 2019.

References II

- [7] Junyi Li, Xiaoxue Cheng, Wayne Xin Zhao, Jian-Yun Nie, and Ji-Rong Wen. Halueval: A large-scale hallucination evaluation benchmark for large language models. *arXiv preprint arXiv:2305.11747*, 2023.
- [8] Stephanie Lin, Jacob Hilton, and Owain Evans. Truthfulqa: Measuring how models mimic human falsehoods. *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics*, pages 3214–3252, 2022.
- [9] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, Peter J Liu, et al. Exploring the limits of transfer learning with a unified text-to-text transformer. *The Journal of Machine Learning Research*, 21(1):5485–5551, 2020.
- [10] Ramprasaath R Selvaraju, Michael Cogswell, Abhishek Das, Ramakrishna Vedantam, Devi Parikh, and Dhruv Batra. Grad-cam: Visual explanations from deep networks via gradient-based localization. In *Proceedings of the IEEE international conference on computer vision*, pages 618–626, 2017.
- [11] James Thorne, Andreas Vlachos, Christos Christodoulopoulos, and Arpit Mittal. Fever: a large-scale dataset for fact extraction and verification. *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 809–819, 2018.

References III

- [12] Naftali Tishby, Fernando C Pereira, and William Bialek. The information bottleneck method. *arXiv preprint physics/0004057*, 2000.
- [13] Aaron van den Oord, Yazhe Li, and Oriol Vinyals. Representation learning with contrastive predictive coding. In *arXiv preprint arXiv:1807.03748*, 2018.
- [14] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in neural information processing systems*, 30, 2017.